

Submission instructions

Submission in pairs unless otherwise authorized

- This notebook contains all the questions. You should follow the instructions below.
- Solutions for both theoretical and practical parts should be provided in this notebook

Moodle submission

You should submit three files:

- IPYNB notebook:
 - All the wet and dry parts, including code, graphs, discussion, etc.
- PDF file:
 - Export the notebook to PDF. Make sure that all the cells are visible.
- Pickle files:
 - As requested in Q3.2.a and Q4.a

All files should be in the following format: "HW1_ID1_ID2.file"

Good Luck!

Question 1 (15pt)

We mentioned in the lecture and tutorial that neural networks, thanks to their non-linear activation functions, can express a much wider range of functions than just linear ones.

In each of the following exercises, you are given the weight matrices, biases and activations of a small neural network. Express the network's output as a simplified, well-known mathematical function and provide a short explanation. There is no need to prove correctness, an intuitive explanation is enough.

The activation function (non-linearity) is denoted as g . The output is given by:

$$s = W_2g(W_1x + b_1) + b_2.$$

In network 4, we also apply the activation to the final output, meaning the network's output is $g(s)$. When not stated otherwise, the bias vector is 0.

Network 1

$$x \in \mathbb{R}. \quad W_1 = \begin{pmatrix} 1 \\ -1 \end{pmatrix}, W_2 = (1 \ 1). \quad g = \text{ReLU}.$$

Solution:

$$s = W_2 \cdot g(W_1 \cdot x + b_1) + b_2$$

Hidden layer:

$$W_1 \cdot x + b_1 = (1, -1)^T \cdot x + (0, 0)^T = (x, -x)$$

Apply ReLU elementwise:

$$g(W_1 \cdot x + b_1) = (\max(0, x), \max(0, -x))$$

Output layer:

$$s = (1, 1) \cdot (\max(0, x), \max(0, -x)) + 0 = \max(0, x) + \max(0, -x)$$

Now check the two cases:

- If $x \geq 0$: $\max(0, x) = x$ and $\max(0, -x) = 0 \Rightarrow s = x$
- If $x < 0$: $\max(0, x) = 0$ and $\max(0, -x) = -x \Rightarrow s = -x$

So in both cases s is the absolute value of x .

Conclusion:

The network implements the function

$$s = |x| \text{ (the absolute value of } x\text{)}.$$

Network 2

$$x \in \mathbb{R}^2. \quad W_1 = \begin{pmatrix} -1 & 1 \\ 1 & 0 \\ -1 & 0 \end{pmatrix}, W_2 = \begin{pmatrix} 1 \\ 1 \\ -1 \end{pmatrix}. \quad g = \text{ReLU}.$$

Network 2 - Solution

1. Hidden layer calculation: $W_1 x$

We first compute the linear transformation into the hidden layer:

$$W_1 x = \begin{pmatrix} -1 & 1 \\ 1 & 0 \\ -1 & 0 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} -1 \cdot x_1 + 1 \cdot x_2 \\ 1 \cdot x_1 + 0 \cdot x_2 \\ -1 \cdot x_1 + 0 \cdot x_2 \end{pmatrix} = \begin{pmatrix} -x_1 + x_2 \\ x_1 \\ -x_1 \end{pmatrix}.$$

Let

$$z = W_1 x = \begin{pmatrix} z_1 \\ z_2 \\ z_3 \end{pmatrix} = \begin{pmatrix} -x_1 + x_2 \\ x_1 \\ -x_1 \end{pmatrix}.$$

2. Hidden layer activation: $h = g(z) = \text{ReLU}(z)$

ReLU is applied elementwise:

$$h = g(z) = \text{ReLU}(z) = \begin{pmatrix} \max(0, z_1) \\ \max(0, z_2) \\ \max(0, z_3) \end{pmatrix} = \begin{pmatrix} \max(0, -x_1 + x_2) \\ \max(0, x_1) \\ \max(0, -x_1) \end{pmatrix}.$$

So the hidden layer output is

$$h = \begin{pmatrix} h_1 \\ h_2 \\ h_3 \end{pmatrix} = \begin{pmatrix} \max(0, -x_1 + x_2) \\ \max(0, x_1) \\ \max(0, -x_1) \end{pmatrix}$$

3. Output layer: $s = W_2^\top h$

Now we compute the scalar output s :

$$s = W_2^\top h = (1 \quad 1 \quad -1) \begin{pmatrix} h_1 \\ h_2 \\ h_3 \end{pmatrix} = 1 \cdot h_1 + 1 \cdot h_2 - 1 \cdot h_3$$

Substitute h_1, h_2, h_3 :

$$s = \max(0, -x_1 + x_2) + \max(0, x_1) - \max(0, -x_1)$$

So far we have

$$s(x_1, x_2) = \max(0, -x_1 + x_2) + \max(0, x_1) - \max(0, -x_1)$$

4. Simplification

Use the identity

$$\max(0, x_1) - \max(0, -x_1) = x_1,$$

which holds for any real x_1 .

Therefore:

$$s = \max(0, -x_1 + x_2) + [\max(0, x_1) - \max(0, -x_1)] = \max(0, -x_1 + x_2) + x_1$$

Note that $-x_1 + x_2 = x_2 - x_1$, so:

$$s = x_1 + \max(0, x_2 - x_1)$$

Now consider two cases:

- If $x_2 \leq x_1$, then $x_2 - x_1 \leq 0$, so $\max(0, x_2 - x_1) = 0$ and

$$s = x_1$$

- If $x_2 > x_1$, then $x_2 - x_1 > 0$, so $\max(0, x_2 - x_1) = x_2 - x_1$ and

$$s = x_1 + (x_2 - x_1) = x_2$$

In both cases, s is the larger of x_1 and x_2 , so:

$$s(x_1, x_2) = \max(x_1, x_2)$$

Final Answer

The network computes the **maximum function** of its two inputs:

$$s(x_1, x_2) = \max(x_1, x_2)$$

Intuition:

The hidden layer constructs $\max(0, x_2 - x_1)$ (active only when $x_2 > x_1$) and also reconstructs x_1 via the combination $\max(0, x_1) - \max(0, -x_1)$. Adding these terms yields the larger of x_1 and x_2 .

Network 3

$$x = \begin{pmatrix} x_1 \\ x_2 \\ y_1 \\ y_2 \end{pmatrix} \in \mathbb{R}^4. \quad W_1 = \begin{pmatrix} 1 & 0 & -1 & 0 \\ -1 & 0 & 1 & 0 \\ 0 & 1 & 0 & -1 \\ 0 & -1 & 0 & 1 \end{pmatrix}, \quad W_2 = (1 \ 1 \ 1 \ 1). \quad g = \text{ReLU}$$



Solution:

Hidden layer:

$$W_1 \cdot x = (x_1 - y_1, -x_1 + y_1, x_2 - y_2, -x_2 + y_2)^T$$

Apply ReLU elementwise:

$$g(W_1 \cdot x) = (\max(0, x_1 - y_1), \max(0, y_1 - x_1), \max(0, x_2 - y_2), \max(0, y_2 - x_2))^T$$

Output:

$$s = W_2 \cdot g(W_1 \cdot x) + b_2$$

$$= \max(0, x_1 - y_1) + \max(0, y_1 - x_1) + \max(0, x_2 - y_2) + \max(0, y_2 - x_2)$$

For any real number a , we have: $\max(0, a) + \max(0, -a) = |a|$.

Using this identity:

$$s = |x_1 - y_1| + |x_2 - y_2|$$

Conclusion:

Network 3 outputs the L1 (Manhattan) distance between the points (x_1, x_2) and (y_1, y_2) :

$$s = |x_1 - y_1| + |x_2 - y_2|.$$

Network 4

Here, the non-linearity is applied also to the output of the second layer. For simplicity, assume $\forall x \geq 10, \text{sigmoid}(x) = 1, \text{sigmoid}(-x) = 0$.

$$x \in \{0, 1\}^2 \quad (\text{binary vector}).$$

$$W_1 = \begin{pmatrix} 20 & 20 \\ -20 & -20 \end{pmatrix}, b_1 = \begin{pmatrix} -10 \\ 30 \end{pmatrix}, W_2 = (20 \ 20), b_2 = -30. \quad g = \text{sigmoid}.$$

Network 4 - Solution

1. Hidden layer pre-activation: $z = W_1 x + b_1$

Let $x = (x_1, x_2)^\top$, Then:

$$z = \begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = W_1 x + b_1 = \begin{pmatrix} 20 & 20 \\ -20 & -20 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} -10 \\ 30 \end{pmatrix}$$

Compute each component:

$$z_1 = 20x_1 + 20x_2 - 10 = 20(x_1 + x_2) - 10$$

$$z_2 = -20x_1 - 20x_2 + 30 = 30 - 20(x_1 + x_2)$$

Since $x_1, x_2 \in \{0, 1\}$, the sum $x_1 + x_2$ can be 0, 1, or 2

2. Hidden layer activation: $h = g(z)$

We apply the (approximate) sigmoid elementwise:

$$h = \begin{pmatrix} h_1 \\ h_2 \end{pmatrix} = \begin{pmatrix} \text{sigmoid}(z_1) \\ \text{sigmoid}(z_2) \end{pmatrix}$$

We now evaluate z_1, z_2, h_1, h_2 for all possible binary inputs.

Case 1: $x = (0, 0)$

Then $x_1 + x_2 = 0$:

$$z_1 = 20 \cdot 0 - 10 = -10 \Rightarrow h_1 = \text{sigmoid}(-10) \approx 0$$

$$z_2 = 30 - 20 \cdot 0 = 30 \Rightarrow h_2 = \text{sigmoid}(30) \approx 1$$

So:

$$h = (0, 1)$$

Case 2: $x = (1, 0)$

Then $x_1 + x_2 = 1$:

$$z_1 = 20 \cdot 1 - 10 = 10 \Rightarrow h_1 = \text{sigmoid}(10) \approx 1$$

$$z_2 = 30 - 20 \cdot 1 = 10 \Rightarrow h_2 = \text{sigmoid}(10) \approx 1$$

So:

$$h = (1, 1)$$

Case 3: $x = (0, 1)$

Symmetric to the previous case ($x_1 + x_2 = 1$), so again:

$$h = (1, 1)$$

Case 4: $x = (1, 1)$

Then $x_1 + x_2 = 2$:

$$z_1 = 20 \cdot 2 - 10 = 30 \Rightarrow h_1 = \text{sigmoid}(30) \approx 1$$

$$z_2 = 30 - 20 \cdot 2 = -10 \Rightarrow h_2 = \text{sigmoid}(-10) \approx 0$$

So:

$$h = (1, 0)$$

We can summarize the hidden layer as:

- $x = (0, 0) \Rightarrow h = (0, 1)$
- $x = (1, 0) \Rightarrow h = (1, 1)$
- $x = (0, 1) \Rightarrow h = (1, 1)$
- $x = (1, 1) \Rightarrow h = (1, 0)$

3. Output pre-activation: $s = W_2 h + b_2$

We have:

$$s = W_2 h + b_2 = (20 \quad 20) \begin{pmatrix} h_1 \\ h_2 \end{pmatrix} - 30 = 20h_1 + 20h_2 - 30$$

Now plug in each h :

- For $x = (0, 0)$: $h = (0, 1)$

$$s = 20 \cdot 0 + 20 \cdot 1 - 30 = -10$$

- For $x = (1, 0)$ or $x = (0, 1)$: $h = (1, 1)$

$$s = 20 \cdot 1 + 20 \cdot 1 - 30 = 10$$

- For $x = (1, 1)$: $h = (1, 0)$

$$s = 20 \cdot 1 + 20 \cdot 0 - 30 = -10$$

So:

- $x = (0, 0) \Rightarrow s = -10$

- $x = (1, 0) \Rightarrow s = 10$
- $x = (0, 1) \Rightarrow s = 10$
- $x = (1, 1) \Rightarrow s = -10$

4. Final output: $y = g(s) = \text{sigmoid}(s)$

Using the same approximation:

- If $s \geq 10$, then $\text{sigmoid}(s) \approx 1$
- If $s \leq -10$, then $\text{sigmoid}(s) \approx 0$

Thus:

- For $x = (0, 0)$: $s = -10 \Rightarrow y \approx 0$
- For $x = (1, 0)$: $s = 10 \Rightarrow y \approx 1$
- For $x = (0, 1)$: $s = 10 \Rightarrow y \approx 1$
- For $x = (1, 1)$: $s = -10 \Rightarrow y \approx 0$

This is exactly the truth table of the **XOR** function:

x_1	x_2	y
0	0	0
1	0	1
0	1	1
1	1	0

So the network computes:

$$y(x_1, x_2) = x_1 \oplus x_2$$

(the XOR of the two input bits).

Intuitive Explanation

- The first hidden neuron (h_1) turns **on** when **at least one** of x_1, x_2 is 1 (since $z_1 = 20(x_1 + x_2) - 10 \geq 10$ whenever $x_1 + x_2 \geq 1$).
 - So $h_1 \approx \mathbf{1}\{x_1 + x_2 \geq 1\}$
- The second hidden neuron (h_2) turns **on** when **not both** inputs are 1 (since $z_2 = 30 - 20(x_1 + x_2) \geq 10$ when $x_1 + x_2 \leq 1$).
 - So $h_2 \approx \mathbf{1}\{x_1 + x_2 \leq 1\}$

Both h_1 and h_2 are 1 **only when exactly one** of x_1, x_2 is 1, i.e. when $x_1 \oplus x_2 = 1$. The second layer then detects this pattern and outputs 1 only in that case.

Therefore, the network implements the **XOR** function:

$$y(x_1, x_2) = x_1 \oplus x_2$$

Question 2

We recommend checking the hidden slides in tutorial 1 before approaching these exercises.

I. Softmax Derivative (5pt)

Derive the gradients of the softmax function and demonstrate how the expression can be reformulated solely by using the softmax function, i.e., in some expression where only $\text{softmax}(x)$, but not x , is present. Recall that the softmax function is defined as follows:

$$\text{softmax}(x)_i = \frac{e^{x_i}}{\sum_{j=1}^N e^{x_j}}$$

I. Softmax Derivative - Answer:

$$\frac{\partial \text{softmax}(x)_i}{\partial x_k} = \dots$$

Derivation

Define

$$s_i = \text{softmax}(x)_i, \quad S = \sum_{j=1}^N e^{x_j},$$

so

$$s_i = \frac{e^{x_i}}{S}.$$

We want

$$\frac{\partial s_i}{\partial x_k}.$$

Using the quotient rule:

$$\frac{\partial s_i}{\partial x_k} = \frac{S \cdot \frac{\partial}{\partial x_k} e^{x_i} - e^{x_i} \cdot \frac{\partial}{\partial x_k} S}{S^2}.$$

1. Derivative of the numerator:

- If $i = k$: $\left(\frac{\partial}{\partial x_k} e^{x_i} = e^{x_i}\right)$
- If $i \neq k$: $\left(\frac{\partial}{\partial x_k} e^{x_i} = 0\right)$

Compactly:

$$\frac{\partial}{\partial x_k} e^{x_i} = e^{x_i} \delta_{ik}.$$

2. Derivative of the denominator:

$$\frac{\partial S}{\partial x_k} = \frac{\partial}{\partial x_k} \sum_{j=1}^N e^{x_j} = e^{x_k}.$$

Substitute into the quotient rule:

$$\frac{\partial s_i}{\partial x_k} = \frac{S \cdot e^{x_i} \delta_{ik} - e^{x_i} \cdot e^{x_k}}{S^2}.$$

Factor $\cancel{(\frac{e^{x_i}}{S})}$:

$$\frac{\partial s_i}{\partial x_k} = \frac{e^{x_i}}{S} \left(\delta_{ik} - \frac{e^{x_k}}{S} \right).$$

But

$$\frac{e^{x_i}}{S} = s_i = \text{softmax}(x)_i, \quad \frac{e^{x_k}}{S} = s_k = \text{softmax}(x)_k,$$

so

$$\frac{\partial s_i}{\partial x_k} = \text{softmax}(x)_i (\delta_{ik} - \text{softmax}(x)_k).$$

Equivalently, in two-case form:

- If $(i = k)$:

$$\frac{\partial s_i}{\partial x_i} = s_i (1 - s_i)$$

- If $\cancel{(i \neq k)}$:

$$\frac{\partial s_i}{\partial x_k} = -s_i s_k$$

II. Cross-Entropy Gradient (5pt)

Derive the gradient of cross-entropy loss with regard to the inputs of a softmax function. i.e., find the gradients with respect to the softmax input vector θ , when the prediction is denoted by $\hat{y} = \text{softmax}(\theta)$. Remember the cross entropy function is:

$$CE(y, \hat{y}) = - \sum_i y_i \log(\hat{y}_i)$$

where y is the one-hot label vector, and \hat{y} is the predicted probability vector for all classes.

II. Cross-Entropy Gradient - Answer

1. Derivative of CE w.r.t. \hat{y}

Start from:

$$L = CE(y, \hat{y}) = - \sum_i y_i \log(\hat{y}_i)$$

For each component \hat{y}_i :

$$\frac{\partial L}{\partial \hat{y}_i} = - \frac{y_i}{\hat{y}_i}$$

2. Derivative of \hat{y} w.r.t. θ (softmax Jacobian)

Softmax:

$$\hat{y}_i = \frac{e^{\theta_i}}{\sum_k e^{\theta_k}}$$

The derivative of \hat{y}_i with respect to θ_j is:

- if $i = j$:

$$\frac{\partial \hat{y}_i}{\partial \theta_j} = \hat{y}_i (1 - \hat{y}_j)$$

- if $i \neq j$:

$$\frac{\partial \hat{y}_i}{\partial \theta_j} = -\hat{y}_i \hat{y}_j$$

In compact form:

$$\frac{\partial \hat{y}_i}{\partial \theta_j} = \hat{y}_i (\delta_{ij} - \hat{y}_j),$$

where δ_{ij} is 1 if $i = j$ and 0 otherwise.

3. Chain rule: $\frac{\partial L}{\partial \theta_j}$

Use the chain rule:

$$\frac{\partial L}{\partial \theta_j} = \sum_i \frac{\partial L}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial \theta_j}$$

Substitute the expressions:

$$\frac{\partial L}{\partial \theta_j} = \sum_i \left(-\frac{y_i}{\hat{y}_i} \right) \hat{y}_i (\delta_{ij} - \hat{y}_j)$$

The \hat{y}_i cancels:

$$\frac{\partial L}{\partial \theta_j} = \sum_i (-y_i(\delta_{ij} - \hat{y}_j))$$

Split the sum:

$$\frac{\partial L}{\partial \theta_j} = \sum_i (-y_i \delta_{ij}) + \sum_i (y_i \hat{y}_j)$$

- First term: only $i = j$ survives:

$$\sum_i (-y_i \delta_{ij}) = -y_j$$

- Second term: \hat{y}_j does not depend on i :

$$\sum_i (y_i \hat{y}_j) = \hat{y}_j \sum_i y_i$$

Because y is one-hot, $\sum_i y_i = 1$, so:

$$\sum_i (y_i \hat{y}_j) = \hat{y}_j$$

Therefore:

$$\frac{\partial L}{\partial \theta_j} = -y_j + \hat{y}_j = \hat{y}_j - y_j$$

4. Final vector form

Stacking all components j into a vector:

$$\nabla_{\theta} L = \hat{y} - y$$

Answer

The gradient of the cross-entropy loss with respect to the softmax input vector θ is:

$$\frac{\partial CE(y, \text{softmax}(\theta))}{\partial \theta} = \hat{y} - y$$

where $\hat{y} = \text{softmax}(\theta)$.

Question 3

I. Derivative Of Activation Functions (5pt)

The following cell contains an implementation of some activation functions. Implement the corresponding derivatives.

```
In [14]: import torch

def sigmoid(x):
    return 1 / (1 + torch.exp(-x))

def tanh(x):
    return torch.div(torch.exp(x) - torch.exp(-x), torch.exp(x) + torch.exp(-x))

def softmax(x):
    exp_x = torch.exp(x.T - torch.max(x, dim=-1).values).T # Subtracting max(x)
    return exp_x / exp_x.sum(dim=-1, keepdim=True)
```

```
In [15]: def d_sigmoid(x):
    s = sigmoid(x)
    return s * (1 - s)

def d_tanh(x):
    t = tanh(x)
    return 1 - t**2

def d_softmax(x):
    s = softmax(x)
    s_diag = torch.diag_embed(s)
    outer = s.unsqueeze(2) * s.unsqueeze(1)
    return s_diag - outer
```

II. Train a Fully Connected network on MNIST (30pt)

In the following exercise, you will create a classifier for the MNIST dataset. You should write your own training and evaluation code and meet the following constraints:

- You are only allowed to use torch tensor manipulations.
- You are NOT allowed to use:
 - Auto-differentiation - backward()
 - Built-in loss functions
 - Built-in activations
 - Built-in optimization
 - Built-in layers (torch.nn)

a. The required classifier class is defined.

You should implement the forward and backward passes of the model. Train the model and plot the model's accuracy and loss (both on train and test sets) as a function of the epochs. You should save the model's weights and biases. Change the student_ids to yours.

In this section, you **must** use the "set_seed" function with the given seed and **sigmoid** as an activation function.

```
In [16]: import torch
import torchvision
from torch.utils.data import DataLoader

import os
import matplotlib.pyplot as plt

# Constants
SEED = 42
EPOCHS = 16
BATCH_SIZE = 32
NUM_OF_CLASSES = 10

# Setting seed
def set_seed(seed):
    torch.manual_seed(seed)
    torch.cuda.manual_seed(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False
    os.environ["PYTHONHASHSEED"] = str(seed)

# Transformation for the data
transform = torchvision.transforms.Compose(
    [torchvision.transforms.ToTensor(),
     torch.flatten])

# Cross-Entropy Loss implementation
def one_hot(y, num_of_classes=10):
    hot = torch.zeros((y.size()[0], num_of_classes))
    hot[torch.arange(y.size()[0]), y] = 1
    return hot

def cross_entropy(y, y_hat):
    return -torch.sum(one_hot(y) * torch.log(y_hat)) / y.size()[0]
```



```
In [17]: # Create dataloaders
train_dataset = torchvision.datasets.MNIST(root='./data', train=True,
                                             download=True, transform=transform)
train_dataloader = torch.utils.data.DataLoader(train_dataset, batch_size=BATCH_SIZ

test_dataset = torchvision.datasets.MNIST(root='./data', train=False,
                                         download=True, transform=transform)
test_dataloader = torch.utils.data.DataLoader(test_dataset, batch_size=BATCH_SIZ
```



```
In [18]: class FullyConnectedNetwork:
    def __init__(self, input_size, output_size, hidden_size1, activation_func,
                 # parameters
                 self.input_size = input_size
                 self.output_size = output_size
                 self.hidden_size1 = hidden_size1

                 # activation function
                 self.activation_func = activation_func

                 # weights
                 self.W1 = torch.randn(self.input_size, self.hidden_size1)
```

```

        self.b1 = torch.zeros(self.hidden_size1)

        self.W2 = torch.randn(self.hidden_size1, self.output_size)
        self.b2 = torch.zeros(self.output_size)

        self.lr = lr

    def forward(self, x):
        self.x = x

        # Hidden layer
        self.z1 = torch.matmul(x, self.W1) + self.b1
        self.h1 = self.activation_func(self.z1)

        # Output layer
        self.z2 = torch.matmul(self.h1, self.W2) + self.b2

        # Probabilities
        self.y_hat = softmax(self.z2)
        return self.y_hat

    def backward(self, x, y, y_hat):
        batch_size = y.size(0)

        # One-hot Labels
        y_one_hot = one_hot(y, num_of_classes=self.output_size)

        # Simplified gradient for softmax + cross-entropy (proved in question 2)
        dl_dz2 = (y_hat - y_one_hot) / batch_size

        # Gradients for W2, b2
        grad_W2 = torch.matmul(self.h1.T, dl_dz2)
        grad_b2 = dl_dz2.sum(dim=0)

        # Backprop to hidden layer
        dl_dh1 = torch.matmul(dl_dz2, self.W2.T)

        # Use d_sigmoid on pre-activation
        dl_dz1 = dl_dh1 * d_sigmoid(self.z1)

        # Gradients for W1, b1
        grad_W1 = torch.matmul(x.T, dl_dz1)
        grad_b1 = dl_dz1.sum(dim=0)

        # SGD update
        self.W1 -= self.lr * grad_W1
        self.b1 -= self.lr * grad_b1
        self.W2 -= self.lr * grad_W2
        self.b2 -= self.lr * grad_b2

```

In [19]:

```
set_seed(SEED)
model = FullyConnectedNetwork(784, 10, 128, sigmoid, lr=0.01)
```

In [20]:

```
# TODO
# Write a training loop
for epoch in range(EPOCHS):
    for x_batch, y_batch in train_dataloader:
        y_hat_batch = model.forward(x_batch)
```

```
loss = cross_entropy(y_batch, y_hat_batch)
model.backward(x_batch, y_batch, y_hat_batch)
```

In [21]:

```
students_ids = "211644240_318801792"
torch.save({"W1": model.W1, "W2": model.W2, "b1": model.b1, "b2": model.b2}, f"{"
```

b. Train the model with various learning rates (at least 3).

Plot the model's accuracy and loss (both on train and test sets) as a function of the epochs. Discuss the differences in training with different learning rates. Support your answer with plots.

In [22]:

```
# Part b - various LRs
learning_rates = [0.001, 0.01, 0.1]
histories = {}

for lr in learning_rates:
    print(f"\nTraining with learning rate {lr}:")

    set_seed(SEED)
    net = FullyConnectedNetwork(784, 10, 128, sigmoid, lr=lr)

    train_losses_lr = []
    test_losses_lr = []
    train_accs_lr = []
    test_accs_lr = []

    for epoch in range(EPOCHS):
        # Train
        net_train_loss = 0.0
        correct_train = 0
        total_train = 0

        for x_batch, y_batch in train_dataloader:
            y_hat_batch = net.forward(x_batch)
            loss = cross_entropy(y_batch, y_hat_batch)
            net_train_loss += loss.item()

            preds = torch.argmax(y_hat_batch, dim=1)
            correct_train += (preds == y_batch).sum().item()
            total_train += y_batch.size(0)

            net.backward(x_batch, y_batch, y_hat_batch)

        avg_train_loss_lr = net_train_loss / len(train_dataloader)
        train_acc_lr = correct_train / total_train

        train_losses_lr.append(avg_train_loss_lr)
        train_accs_lr.append(train_acc_lr)

        # Evaluate on test set
        net_test_loss = 0.0
        correct_test = 0
        total_test = 0

        for x_test, y_test in test_dataloader:
            y_hat_test = net.forward(x_test)
            loss_test = cross_entropy(y_test, y_hat_test)
```

```
net_test_loss += loss_test.item()

preds_test = torch.argmax(y_hat_test, dim=1)
correct_test += (preds_test == y_test).sum().item()
total_test += y_test.size(0)

avg_test_loss_lr = net_test_loss / len(test_dataloader)
test_acc_lr = correct_test / total_test

test_losses_lr.append(avg_test_loss_lr)
test_accs_lr.append(test_acc_lr)

print(f"[lr={lr}] Epoch {epoch+1}/{EPOCHS} "
      f"- Train Loss: {avg_train_loss_lr:.4f}, Train Acc: {train_acc_lr:.4f}"
      f"- Test Loss: {avg_test_loss_lr:.4f}, Test Acc: {test_acc_lr:.4f}")

histories[lr] = {"train_losses": train_losses_lr,
                 "test_losses": test_losses_lr,
                 "train_accs": train_accs_lr,
                 "test_accs": test_accs_lr,}
```

Training with learning rate 0.001:

```
[lr=0.001] Epoch 1/16 - Train Loss: 8.0070, Train Acc: 0.1355 - Test Loss: 5.989  
1, Test Acc: 0.1709  
[lr=0.001] Epoch 2/16 - Train Loss: 5.2537, Train Acc: 0.1931 - Test Loss: 4.682  
1, Test Acc: 0.2262  
[lr=0.001] Epoch 3/16 - Train Loss: 4.3004, Train Acc: 0.2521 - Test Loss: 3.892  
1, Test Acc: 0.2872  
[lr=0.001] Epoch 4/16 - Train Loss: 3.6346, Train Acc: 0.3108 - Test Loss: 3.318  
1, Test Acc: 0.3497  
[lr=0.001] Epoch 5/16 - Train Loss: 3.1449, Train Acc: 0.3647 - Test Loss: 2.893  
1, Test Acc: 0.3986  
[lr=0.001] Epoch 6/16 - Train Loss: 2.7770, Train Acc: 0.4114 - Test Loss: 2.570  
4, Test Acc: 0.4423  
[lr=0.001] Epoch 7/16 - Train Loss: 2.4942, Train Acc: 0.4517 - Test Loss: 2.319  
6, Test Acc: 0.4763  
[lr=0.001] Epoch 8/16 - Train Loss: 2.2722, Train Acc: 0.4860 - Test Loss: 2.120  
6, Test Acc: 0.5083  
[lr=0.001] Epoch 9/16 - Train Loss: 2.0944, Train Acc: 0.5161 - Test Loss: 1.959  
4, Test Acc: 0.5373  
[lr=0.001] Epoch 10/16 - Train Loss: 1.9492, Train Acc: 0.5415 - Test Loss: 1.826  
6, Test Acc: 0.5622  
[lr=0.001] Epoch 11/16 - Train Loss: 1.8287, Train Acc: 0.5620 - Test Loss: 1.715  
3, Test Acc: 0.5853  
[lr=0.001] Epoch 12/16 - Train Loss: 1.7269, Train Acc: 0.5813 - Test Loss: 1.620  
7, Test Acc: 0.6024  
[lr=0.001] Epoch 13/16 - Train Loss: 1.6398, Train Acc: 0.5990 - Test Loss: 1.539  
3, Test Acc: 0.6183  
[lr=0.001] Epoch 14/16 - Train Loss: 1.5644, Train Acc: 0.6132 - Test Loss: 1.468  
4, Test Acc: 0.6327  
[lr=0.001] Epoch 15/16 - Train Loss: 1.4983, Train Acc: 0.6266 - Test Loss: 1.406  
1, Test Acc: 0.6438  
[lr=0.001] Epoch 16/16 - Train Loss: 1.4398, Train Acc: 0.6386 - Test Loss: 1.350  
9, Test Acc: 0.6544
```

Training with learning rate 0.01:

```
[lr=0.01] Epoch 1/16 - Train Loss: 3.5984, Train Acc: 0.3653 - Test Loss: 1.8211,  
Test Acc: 0.5635  
[lr=0.01] Epoch 2/16 - Train Loss: 1.5013, Train Acc: 0.6265 - Test Loss: 1.1778,  
Test Acc: 0.6916  
[lr=0.01] Epoch 3/16 - Train Loss: 1.1075, Train Acc: 0.7099 - Test Loss: 0.9392,  
Test Acc: 0.7459  
[lr=0.01] Epoch 4/16 - Train Loss: 0.9252, Train Acc: 0.7512 - Test Loss: 0.8109,  
Test Acc: 0.7759  
[lr=0.01] Epoch 5/16 - Train Loss: 0.8157, Train Acc: 0.7775 - Test Loss: 0.7291,  
Test Acc: 0.7961  
[lr=0.01] Epoch 6/16 - Train Loss: 0.7406, Train Acc: 0.7966 - Test Loss: 0.6715,  
Test Acc: 0.8104  
[lr=0.01] Epoch 7/16 - Train Loss: 0.6848, Train Acc: 0.8105 - Test Loss: 0.6282,  
Test Acc: 0.8237  
[lr=0.01] Epoch 8/16 - Train Loss: 0.6411, Train Acc: 0.8217 - Test Loss: 0.5939,  
Test Acc: 0.8314  
[lr=0.01] Epoch 9/16 - Train Loss: 0.6056, Train Acc: 0.8304 - Test Loss: 0.5659,  
Test Acc: 0.8381  
[lr=0.01] Epoch 10/16 - Train Loss: 0.5760, Train Acc: 0.8381 - Test Loss: 0.542  
4, Test Acc: 0.8443  
[lr=0.01] Epoch 11/16 - Train Loss: 0.5508, Train Acc: 0.8447 - Test Loss: 0.522  
2, Test Acc: 0.8494  
[lr=0.01] Epoch 12/16 - Train Loss: 0.5289, Train Acc: 0.8505 - Test Loss: 0.504  
7, Test Acc: 0.8545  
[lr=0.01] Epoch 13/16 - Train Loss: 0.5097, Train Acc: 0.8555 - Test Loss: 0.489
```

```

2, Test Acc: 0.8591
[lr=0.01] Epoch 14/16 - Train Loss: 0.4926, Train Acc: 0.8602 - Test Loss: 0.475
4, Test Acc: 0.8644
[lr=0.01] Epoch 15/16 - Train Loss: 0.4773, Train Acc: 0.8640 - Test Loss: 0.463
1, Test Acc: 0.8674
[lr=0.01] Epoch 16/16 - Train Loss: 0.4634, Train Acc: 0.8676 - Test Loss: 0.451
9, Test Acc: 0.8701

Training with learning rate 0.1:
[lr=0.1] Epoch 1/16 - Train Loss: 1.1455, Train Acc: 0.7284 - Test Loss: 0.5417,
Test Acc: 0.8481
[lr=0.1] Epoch 2/16 - Train Loss: 0.4878, Train Acc: 0.8622 - Test Loss: 0.4182,
Test Acc: 0.8791
[lr=0.1] Epoch 3/16 - Train Loss: 0.3866, Train Acc: 0.8883 - Test Loss: 0.3616,
Test Acc: 0.8926
[lr=0.1] Epoch 4/16 - Train Loss: 0.3329, Train Acc: 0.9031 - Test Loss: 0.3267,
Test Acc: 0.9014
[lr=0.1] Epoch 5/16 - Train Loss: 0.2977, Train Acc: 0.9129 - Test Loss: 0.3025,
Test Acc: 0.9087
[lr=0.1] Epoch 6/16 - Train Loss: 0.2723, Train Acc: 0.9203 - Test Loss: 0.2845,
Test Acc: 0.9134
[lr=0.1] Epoch 7/16 - Train Loss: 0.2527, Train Acc: 0.9254 - Test Loss: 0.2705,
Test Acc: 0.9174
[lr=0.1] Epoch 8/16 - Train Loss: 0.2369, Train Acc: 0.9300 - Test Loss: 0.2593,
Test Acc: 0.9216
[lr=0.1] Epoch 9/16 - Train Loss: 0.2237, Train Acc: 0.9341 - Test Loss: 0.2499,
Test Acc: 0.9239
[lr=0.1] Epoch 10/16 - Train Loss: 0.2124, Train Acc: 0.9375 - Test Loss: 0.2420,
Test Acc: 0.9276
[lr=0.1] Epoch 11/16 - Train Loss: 0.2026, Train Acc: 0.9406 - Test Loss: 0.2353,
Test Acc: 0.9300
[lr=0.1] Epoch 12/16 - Train Loss: 0.1939, Train Acc: 0.9434 - Test Loss: 0.2294,
Test Acc: 0.9321
[lr=0.1] Epoch 13/16 - Train Loss: 0.1862, Train Acc: 0.9457 - Test Loss: 0.2243,
Test Acc: 0.9334
[lr=0.1] Epoch 14/16 - Train Loss: 0.1792, Train Acc: 0.9480 - Test Loss: 0.2197,
Test Acc: 0.9342
[lr=0.1] Epoch 15/16 - Train Loss: 0.1728, Train Acc: 0.9497 - Test Loss: 0.2156,
Test Acc: 0.9353
[lr=0.1] Epoch 16/16 - Train Loss: 0.1669, Train Acc: 0.9516 - Test Loss: 0.2119,
Test Acc: 0.9366

```

Question 4

We recommend using GPUs to solve this section. You can use free GPUs provided by Google Colab or by Kaggle (instructions on how to connect are in Moodle). Also, we recommend going over all tutorial notebooks and using dropout, normalization and Adam for better learning. Solving this on your personal computer will be significantly slower than using cloud computation.

I. Implement and Train a CNN (30pt)

Whenever Ben sees a cat, his first instinct is to run up and pet them. However, after a recent incident he realized not all types of cats are very fond of that type of attention.

Train a classifier to tell the difference between different types of big cats, and help Ben from putting himself in too much danger.

Your code should meet the following constraints:

1. Your classifier must be CNN based.
2. You are not allowed to use any pre-trained model.

Stages

1. Perform a short EDA (Exploratory Data Analysis).
2. Train the model and plot its accuracy and loss (for both the training and validation sets) as a function of the epochs. Display the overall number of parameters in the model. To see the number of parameters in your model, run `sum(p.numel() for p in model.parameters())`.
3. Report the test set accuracy.

Your data is in `hw1_data/big_cats`.

You can use the provided code for preprocessing the data. You can also define a custom dataset (as in tutorial 3) or use `torchvision.datasets.ImageFolder` if you prefer.

Your training function must call the `set_seed` command for reproducibility.

Submission

In addition to the code in the notebook, you should submit:

1. a `.py` file containing your model class.
2. a `.pk1` file containing the weight of your model.

Grading

- 5 points for EDA.
- 15 points based on reaching minimum test accuracy (77% accuracy required to get all 15 points, 70% accuracy to get 10 points).
- 10 points competitive: try to balance the tradeoff between number of parameters in your model and model's test accuracy. The lower the # of parameters and the higher the test accuracy, the higher the final score.

```
In [26]: import os
import copy
import random
import numpy as np
import torch
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms
from torchvision.datasets import ImageFolder
import torch.nn as nn
import torch.optim as optim

import matplotlib.pyplot as plt
from PIL import Image
```

```

def set_seed(seed: int):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    if torch.cuda.is_available():
        torch.cuda.manual_seed_all(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

set_seed(42)

class BigCatDataset(Dataset):
    def __init__(self, root_dir, transform=None):
        """
        root_dir: path to 'train', 'valid' or 'test' folder.
                  Inside it: one subfolder per class with images.
        """
        self.root_dir = root_dir
        self.transform = transform

        self.classes = sorted([
            d for d in os.listdir(root_dir)
            if os.path.isdir(os.path.join(root_dir, d))
        ])

        # build list of (path, label)
        self.samples = []
        for class_idx, class_name in enumerate(self.classes):
            class_dir = os.path.join(root_dir, class_name)
            for fname in os.listdir(class_dir):
                if fname.lower().endswith((".jpg", ".jpeg", ".png")):
                    self.samples.append((os.path.join(class_dir, fname), class_i

def __len__(self):
    return len(self.samples)

def __getitem__(self, index):
    img_path, label = self.samples[index]
    img = Image.open(img_path).convert("RGB")
    if self.transform:
        img = self.transform(img)
    return img, label

def get_class(self, idx):
    return self.classes[idx]

# =====
# Transforms
# =====
train_transform = transforms.Compose([
    transforms.Resize((72, 72)),
    transforms.RandomResizedCrop(64, scale=(0.8, 1.0), ratio=(0.9, 1.1)),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomRotation(degrees=10),
    transforms.ColorJitter(brightness=0.15, contrast=0.15, saturation=0.15, hue=
    transforms.ToTensor(),
])

```

```

        transforms.Normalize((0.5, 0.5, 0.5),
                             (0.5, 0.5, 0.5))
    ])

val_test_transform = transforms.Compose([
    transforms.Resize((64, 64)),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5),
                      (0.5, 0.5, 0.5))
])

# =====
# EDA Helper Functions
# =====

def count_images(root):
    per_class = {}
    for cls in os.listdir(root):
        cls_path = os.path.join(root, cls)
        if os.path.isdir(cls_path):
            per_class[cls] = len(os.listdir(cls_path))
    return per_class, sum(per_class.values())


def show_random_samples(dataset, n=6):
    idxs = random.sample(range(len(dataset)), n)
    plt.figure(figsize=(n * 2.2, 3))
    for i, idx in enumerate(idxs):
        img, label = dataset[idx]
        plt.subplot(1, n, i + 1)

        if isinstance(img, torch.Tensor):
            img = img.permute(1, 2, 0).numpy()
            img = (img * 0.5) + 0.5 # roughly unnormalize
        else:
            img = np.array(img) # makes sure PIL → numpy

        plt.imshow(img)
        plt.title(dataset.classes[label])
        plt.axis("off")

    plt.tight_layout()
    plt.show()

class BigCatsCNN(nn.Module):
    """
    Deeper CNN with controlled parameter count (< 500k):

    - 4 conv blocks, each with 2 conv layers:
        3 → 24 → 48 → 96 → 128 channels
    - Each conv: 3x3, padding=1, followed by BN + ReLU
    - MaxPool2d(2) after each block to downsample
    - AdaptiveAvgPool2d((1, 1)) + small classifier
    """
    def __init__(self, num_classes=10):
        super(BigCatsCNN, self).__init__()

        def conv_block(in_ch, out_ch, num_convs=2):
            layers = []

```

```

    for i in range(num_convs):
        ch_in = in_ch if i == 0 else out_ch
        layers += [
            nn.Conv2d(ch_in, out_ch, kernel_size=3, padding=1, bias=False),
            nn.BatchNorm2d(out_ch),
            nn.ReLU(inplace=True),
        ]
    layers.append(nn.MaxPool2d(2)) # downsample by 2
    return nn.Sequential(*layers)

    # Input: (B, 3, 64, 64)
    self.features = nn.Sequential(
        conv_block(3, 24),      # 64 -> 32
        conv_block(24, 48),     # 32 -> 16
        conv_block(48, 96),     # 16 -> 8
        conv_block(96, 128),    # 8 -> 4
        nn.AdaptiveAvgPool2d((1, 1)) # -> (B, 128, 1, 1)
    )

    self.classifier = nn.Sequential(
        nn.Flatten(),           # (B, 128)
        nn.Linear(128, 128),
        nn.ReLU(True),
        nn.Dropout(p=0.5),
        nn.Linear(128, num_classes),
        nn.LogSoftmax(dim=1),
    )

def forward(self, x):
    x = self.features(x)
    x = self.classifier(x)
    return x

# Root data directory
DATA_DIR = "big_cats"
train_dir = os.path.join(DATA_DIR, "train")
valid_dir = os.path.join(DATA_DIR, "valid")
test_dir = os.path.join(DATA_DIR, "test")

print("===== Exploratory Data Analysis =====\n")
train_tmp = ImageFolder(train_dir)
classes = train_tmp.classes
print("==== Classes of Big Cats ====\n")
print(classes)

print("\n==== Count in Split and Classes ====")
for name, d in [("Train", train_dir), ("Valid", valid_dir), ("Test", test_dir)]:
    per_class, total = count_images(d)
    print(f"\n{name} - {total} images total")
    for cls, n in per_class.items():
        print(f"  {cls}: {n}")

print("\n==== Show Samples (Train) ====\n")
sample_ds = ImageFolder(train_dir)
random.seed(42)
show_random_samples(sample_ds, n=6)

print("\n==== Image Size Distribution (Train) ====\n")
sizes = []

```

```

for cls in os.listdir(train_dir):
    cls_path = os.path.join(train_dir, cls)
    if os.path.isdir(cls_path):
        for img_name in os.listdir(cls_path):
            img_path = os.path.join(cls_path, img_name)
            try:
                with Image.open(img_path) as img:
                    sizes.append(img.size)
            except Exception:
                pass

sizes = np.array(sizes)
widths = sizes[:, 0]
heights = sizes[:, 1]

print("Image size statistics (train set):\n")
print(f"Mean      : width {widths.mean():.1f}, height {heights.mean():.1f}")
print(f"Median    : width {np.median(widths):.1f}, height {np.median(heights):.1f}")
print(f"Std dev   : width {widths.std():.1f}, height {heights.std():.1f}")

# ----- Datasets & DataLoaders -----
train_dataset = BigCatDataset(train_dir, transform=train_transform)
val_dataset = BigCatDataset(valid_dir, transform=val_test_transform)
test_dataset = BigCatDataset(test_dir, transform=val_test_transform)

BATCH_SIZE = 32

train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False)

num_classes = len(train_dataset.classes)
print("\nClasses:", train_dataset.classes)
print("Train size:", len(train_dataset),
      "Val size:", len(val_dataset),
      "Test size:", len(test_dataset))

# ---- Compute class weights (smoothed inverse frequency) ----
labels_list = [label for _, label in train_dataset.samples]
class_counts = np.bincount(labels_list, minlength=num_classes)
print("\nClass counts:", class_counts)

# Base inverse-frequency weights
inv_freq = 1.0 / class_counts
inv_freq = inv_freq / inv_freq.mean() # normalize around 1

# Smoothing parameter: 1.0 = original, 0.5 = milder, 0.0 = no weighting
ALPHA = 0.5
class_weights = inv_freq ** ALPHA
class_weights = class_weights / class_weights.mean() # normalize again
class_weights = torch.tensor(class_weights, dtype=torch.float32)

device = "cuda" if torch.cuda.is_available() else "cpu"
print(f"\nUsing device: {device}")

model = BigCatsCNN(num_classes=num_classes).to(device)

# Use class-weighted loss
class_weights = class_weights.to(device)
criterion = nn.NLLLoss(weight=class_weights)

```

```

optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-4)
scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=20, gamma=0.3)

num_params = sum(p.numel() for p in model.parameters())
print("Number of parameters:", num_params)

num_epochs = 150

train_loss_hist = []
val_loss_hist = []
train_acc_hist = []
val_acc_hist = []

best_val_acc = 0.0
best_state_dict = copy.deepcopy(model.state_dict())
best_epoch = 0

# ----- Training Loop -----
for epoch in range(num_epochs):
    # ---- Train ----
    model.train()
    running_loss = 0.0
    running_correct = 0
    running_total = 0

    for images, labels in train_loader:
        images = images.to(device)
        labels = labels.to(device)

        outputs = model(images)
        loss = criterion(outputs, labels)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * images.size(0)
        preds = outputs.argmax(dim=1)
        running_correct += (preds == labels).sum().item()
        running_total += labels.size(0)

    epoch_train_loss = running_loss / running_total
    epoch_train_acc = running_correct / running_total

    # ---- Validation ----
    model.eval()
    val_loss = 0.0
    val_correct = 0
    val_total = 0

    with torch.no_grad():
        for images, labels in val_loader:
            images = images.to(device)
            labels = labels.to(device)

            outputs = model(images)
            loss = criterion(outputs, labels)

            val_loss += loss.item() * images.size(0)

```

```

        preds = outputs.argmax(dim=1)
        val_correct += (preds == labels).sum().item()
        val_total += labels.size(0)

        epoch_val_loss = val_loss / val_total
        epoch_val_acc = val_correct / val_total

        train_loss_hist.append(epoch_train_loss)
        val_loss_hist.append(epoch_val_loss)
        train_acc_hist.append(epoch_train_acc)
        val_acc_hist.append(epoch_val_acc)

        scheduler.step()

    # Track best model
    if epoch_val_acc > best_val_acc:
        best_val_acc = epoch_val_acc
        best_state_dict = copy.deepcopy(model.state_dict())
        best_epoch = epoch + 1

    print(
        f"Epoch [{epoch + 1}/{num_epochs}] "
        f"Train Loss: {epoch_train_loss:.4f} Acc: {epoch_train_acc:.3f} | "
        f"Val Loss: {epoch_val_loss:.4f} Acc: {epoch_val_acc:.3f}"
    )

print(f"\nBest validation accuracy: {best_val_acc:.3f} at epoch {best_epoch}")

# Load best weights into the model
model.load_state_dict(best_state_dict)

# ----- Plots -----
epochs = range(1, num_epochs + 1)

plt.figure()
plt.plot(epochs, train_loss_hist, label="Train Loss")
plt.plot(epochs, val_loss_hist, label="Val Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Loss vs Epochs")
plt.legend()
plt.tight_layout()
plt.show()

plt.figure()
plt.plot(epochs, train_acc_hist, label="Train Acc")
plt.plot(epochs, val_acc_hist, label="Val Acc")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.title("Accuracy vs Epochs")
plt.legend()
plt.tight_layout()
plt.show()

# ----- Test Evaluation -----
model.eval()
correct = 0
total = 0

with torch.no_grad():

```

```
for images, labels in test_loader:
    images = images.to(device)
    labels = labels.to(device)

    outputs = model(images)
    preds = outputs.argmax(dim=1)
    correct += (preds == labels).sum().item()
    total += labels.size(0)

test_acc = correct / total
print(f"\nTest Accuracy: {test_acc:.3f}")

# ----- Save best model weights to .pkl -----
weights_path = "HW1_211644240_318801792_Q4.pkl"
torch.save(best_state_dict, weights_path)
print(f"Saved best model weights to: {weights_path}")
```

===== Exploratory Data Analysis =====

==== Classes of Big Cats ====

```
[ 'AFRICAN LEOPARD', 'CARACAL', 'CHEETAH', 'CLOUDED LEOPARD', 'JAGUAR', 'LIONS',
'OCELOT', 'PUMA', 'SNOW LEOPARD', 'TIGER']
```

==== Count in Split and Classes ====

Train – 2339 images total

```
AFRICAN LEOPARD: 236
CARACAL: 236
CHEETAH: 235
CLOUDED LEOPARD: 229
JAGUAR: 238
LIONS: 228
OCELOT: 233
PUMA: 236
SNOW LEOPARD: 231
TIGER: 237
```

Valid – 50 images total

```
AFRICAN LEOPARD: 5
CARACAL: 5
CHEETAH: 5
CLOUDED LEOPARD: 5
JAGUAR: 5
LIONS: 5
OCELOT: 5
PUMA: 5
SNOW LEOPARD: 5
TIGER: 5
```

Test – 50 images total

```
AFRICAN LEOPARD: 5
CARACAL: 5
CHEETAH: 5
CLOUDED LEOPARD: 5
JAGUAR: 5
LIONS: 5
OCELOT: 5
PUMA: 5
SNOW LEOPARD: 5
TIGER: 5
```

==== Show Samples (Train) ====

==== Image Size Distribution (Train) ====

Image size statistics (train set):

```
Mean      : width 224.0, height 224.0
Median    : width 224.0, height 224.0
Std dev   : width 0.0, height 0.0
```

```
Classes: [ 'AFRICAN LEOPARD', 'CARACAL', 'CHEETAH', 'CLOUDED LEOPARD', 'JAGUAR',
'LIONS', 'OCELOT', 'PUMA', 'SNOW LEOPARD', 'TIGER']
```

```
Train size: 2339 Val size: 50 Test size: 50
```

Class counts: [236 236 235 229 238 228 233 236 231 237]

Using device: cpu

Number of parameters: 438386

Epoch [1/150] Train Loss: 2.0338 Acc: 0.243 | Val Loss: 1.7977 Acc: 0.280
Epoch [2/150] Train Loss: 1.7922 Acc: 0.310 | Val Loss: 1.8507 Acc: 0.320
Epoch [3/150] Train Loss: 1.6528 Acc: 0.340 | Val Loss: 2.1888 Acc: 0.260
Epoch [4/150] Train Loss: 1.5827 Acc: 0.362 | Val Loss: 1.3887 Acc: 0.480
Epoch [5/150] Train Loss: 1.5935 Acc: 0.354 | Val Loss: 1.4092 Acc: 0.460
Epoch [6/150] Train Loss: 1.5141 Acc: 0.406 | Val Loss: 1.4304 Acc: 0.420
Epoch [7/150] Train Loss: 1.4591 Acc: 0.420 | Val Loss: 1.4667 Acc: 0.400
Epoch [8/150] Train Loss: 1.4196 Acc: 0.445 | Val Loss: 1.4588 Acc: 0.500
Epoch [9/150] Train Loss: 1.3550 Acc: 0.472 | Val Loss: 1.6558 Acc: 0.420
Epoch [10/150] Train Loss: 1.3145 Acc: 0.485 | Val Loss: 1.4374 Acc: 0.460
Epoch [11/150] Train Loss: 1.2652 Acc: 0.498 | Val Loss: 1.2231 Acc: 0.600
Epoch [12/150] Train Loss: 1.2118 Acc: 0.533 | Val Loss: 1.3187 Acc: 0.440
Epoch [13/150] Train Loss: 1.1410 Acc: 0.555 | Val Loss: 1.0894 Acc: 0.620
Epoch [14/150] Train Loss: 1.1009 Acc: 0.585 | Val Loss: 1.0954 Acc: 0.560
Epoch [15/150] Train Loss: 1.0958 Acc: 0.572 | Val Loss: 1.4515 Acc: 0.500
Epoch [16/150] Train Loss: 1.0070 Acc: 0.614 | Val Loss: 1.3979 Acc: 0.460
Epoch [17/150] Train Loss: 0.9723 Acc: 0.615 | Val Loss: 1.5607 Acc: 0.500
Epoch [18/150] Train Loss: 0.9451 Acc: 0.640 | Val Loss: 1.1721 Acc: 0.600
Epoch [19/150] Train Loss: 0.9156 Acc: 0.664 | Val Loss: 1.9565 Acc: 0.380
Epoch [20/150] Train Loss: 0.8544 Acc: 0.689 | Val Loss: 1.0422 Acc: 0.640
Epoch [21/150] Train Loss: 0.7830 Acc: 0.720 | Val Loss: 0.8544 Acc: 0.680
Epoch [22/150] Train Loss: 0.7065 Acc: 0.746 | Val Loss: 0.6814 Acc: 0.720
Epoch [23/150] Train Loss: 0.7024 Acc: 0.744 | Val Loss: 0.9871 Acc: 0.640
Epoch [24/150] Train Loss: 0.6834 Acc: 0.762 | Val Loss: 0.8146 Acc: 0.640
Epoch [25/150] Train Loss: 0.6765 Acc: 0.760 | Val Loss: 0.8314 Acc: 0.680
Epoch [26/150] Train Loss: 0.6444 Acc: 0.773 | Val Loss: 0.8286 Acc: 0.700
Epoch [27/150] Train Loss: 0.6618 Acc: 0.767 | Val Loss: 0.8438 Acc: 0.700
Epoch [28/150] Train Loss: 0.6116 Acc: 0.779 | Val Loss: 0.7965 Acc: 0.700
Epoch [29/150] Train Loss: 0.6079 Acc: 0.782 | Val Loss: 0.6226 Acc: 0.720
Epoch [30/150] Train Loss: 0.5670 Acc: 0.787 | Val Loss: 0.6836 Acc: 0.740
Epoch [31/150] Train Loss: 0.5595 Acc: 0.807 | Val Loss: 0.6686 Acc: 0.720
Epoch [32/150] Train Loss: 0.6135 Acc: 0.790 | Val Loss: 0.5597 Acc: 0.840
Epoch [33/150] Train Loss: 0.5660 Acc: 0.797 | Val Loss: 0.7856 Acc: 0.700
Epoch [34/150] Train Loss: 0.5348 Acc: 0.808 | Val Loss: 0.9045 Acc: 0.760
Epoch [35/150] Train Loss: 0.5193 Acc: 0.822 | Val Loss: 0.7120 Acc: 0.720
Epoch [36/150] Train Loss: 0.4943 Acc: 0.826 | Val Loss: 0.5579 Acc: 0.780
Epoch [37/150] Train Loss: 0.4951 Acc: 0.825 | Val Loss: 0.8208 Acc: 0.720
Epoch [38/150] Train Loss: 0.5156 Acc: 0.812 | Val Loss: 0.4308 Acc: 0.860
Epoch [39/150] Train Loss: 0.5250 Acc: 0.820 | Val Loss: 0.6809 Acc: 0.700
Epoch [40/150] Train Loss: 0.4986 Acc: 0.819 | Val Loss: 0.6436 Acc: 0.800
Epoch [41/150] Train Loss: 0.4486 Acc: 0.844 | Val Loss: 0.5648 Acc: 0.800
Epoch [42/150] Train Loss: 0.4144 Acc: 0.861 | Val Loss: 0.5508 Acc: 0.760
Epoch [43/150] Train Loss: 0.4195 Acc: 0.854 | Val Loss: 0.5274 Acc: 0.800
Epoch [44/150] Train Loss: 0.4324 Acc: 0.855 | Val Loss: 0.5183 Acc: 0.800
Epoch [45/150] Train Loss: 0.4151 Acc: 0.869 | Val Loss: 0.4982 Acc: 0.840
Epoch [46/150] Train Loss: 0.4291 Acc: 0.847 | Val Loss: 0.4892 Acc: 0.800
Epoch [47/150] Train Loss: 0.4244 Acc: 0.850 | Val Loss: 0.5623 Acc: 0.780
Epoch [48/150] Train Loss: 0.4156 Acc: 0.859 | Val Loss: 0.5163 Acc: 0.780
Epoch [49/150] Train Loss: 0.4044 Acc: 0.858 | Val Loss: 0.4713 Acc: 0.820
Epoch [50/150] Train Loss: 0.3984 Acc: 0.864 | Val Loss: 0.4344 Acc: 0.820
Epoch [51/150] Train Loss: 0.4003 Acc: 0.862 | Val Loss: 0.4650 Acc: 0.820
Epoch [52/150] Train Loss: 0.3889 Acc: 0.873 | Val Loss: 0.4635 Acc: 0.800
Epoch [53/150] Train Loss: 0.3901 Acc: 0.872 | Val Loss: 0.4892 Acc: 0.820
Epoch [54/150] Train Loss: 0.3743 Acc: 0.881 | Val Loss: 0.5289 Acc: 0.820
Epoch [55/150] Train Loss: 0.3879 Acc: 0.861 | Val Loss: 0.4725 Acc: 0.820
Epoch [56/150] Train Loss: 0.3693 Acc: 0.876 | Val Loss: 0.5033 Acc: 0.820

Epoch [57/150] Train Loss: 0.3548 Acc: 0.877	Val Loss: 0.4535 Acc: 0.840
Epoch [58/150] Train Loss: 0.3741 Acc: 0.870	Val Loss: 0.4027 Acc: 0.840
Epoch [59/150] Train Loss: 0.3805 Acc: 0.869	Val Loss: 0.4444 Acc: 0.860
Epoch [60/150] Train Loss: 0.3776 Acc: 0.865	Val Loss: 0.4050 Acc: 0.820
Epoch [61/150] Train Loss: 0.3564 Acc: 0.876	Val Loss: 0.4096 Acc: 0.860
Epoch [62/150] Train Loss: 0.3398 Acc: 0.885	Val Loss: 0.4409 Acc: 0.840
Epoch [63/150] Train Loss: 0.3441 Acc: 0.888	Val Loss: 0.4224 Acc: 0.820
Epoch [64/150] Train Loss: 0.3454 Acc: 0.879	Val Loss: 0.4574 Acc: 0.820
Epoch [65/150] Train Loss: 0.3644 Acc: 0.873	Val Loss: 0.3923 Acc: 0.840
Epoch [66/150] Train Loss: 0.3451 Acc: 0.881	Val Loss: 0.4305 Acc: 0.880
Epoch [67/150] Train Loss: 0.3307 Acc: 0.881	Val Loss: 0.4049 Acc: 0.820
Epoch [68/150] Train Loss: 0.3244 Acc: 0.886	Val Loss: 0.4813 Acc: 0.800
Epoch [69/150] Train Loss: 0.3350 Acc: 0.886	Val Loss: 0.4136 Acc: 0.820
Epoch [70/150] Train Loss: 0.3383 Acc: 0.892	Val Loss: 0.4351 Acc: 0.820
Epoch [71/150] Train Loss: 0.3319 Acc: 0.881	Val Loss: 0.4107 Acc: 0.820
Epoch [72/150] Train Loss: 0.3264 Acc: 0.888	Val Loss: 0.4087 Acc: 0.820
Epoch [73/150] Train Loss: 0.3429 Acc: 0.885	Val Loss: 0.4035 Acc: 0.820
Epoch [74/150] Train Loss: 0.3494 Acc: 0.881	Val Loss: 0.4297 Acc: 0.800
Epoch [75/150] Train Loss: 0.3548 Acc: 0.880	Val Loss: 0.3912 Acc: 0.880
Epoch [76/150] Train Loss: 0.3490 Acc: 0.878	Val Loss: 0.4497 Acc: 0.760
Epoch [77/150] Train Loss: 0.3195 Acc: 0.885	Val Loss: 0.4433 Acc: 0.840
Epoch [78/150] Train Loss: 0.3336 Acc: 0.880	Val Loss: 0.4593 Acc: 0.820
Epoch [79/150] Train Loss: 0.3273 Acc: 0.889	Val Loss: 0.4283 Acc: 0.840
Epoch [80/150] Train Loss: 0.3252 Acc: 0.888	Val Loss: 0.4348 Acc: 0.860
Epoch [81/150] Train Loss: 0.3170 Acc: 0.889	Val Loss: 0.4134 Acc: 0.840
Epoch [82/150] Train Loss: 0.3365 Acc: 0.889	Val Loss: 0.4160 Acc: 0.840
Epoch [83/150] Train Loss: 0.3343 Acc: 0.898	Val Loss: 0.4436 Acc: 0.860
Epoch [84/150] Train Loss: 0.3056 Acc: 0.895	Val Loss: 0.4099 Acc: 0.840
Epoch [85/150] Train Loss: 0.3263 Acc: 0.899	Val Loss: 0.4350 Acc: 0.820
Epoch [86/150] Train Loss: 0.3140 Acc: 0.897	Val Loss: 0.4093 Acc: 0.840
Epoch [87/150] Train Loss: 0.3238 Acc: 0.891	Val Loss: 0.4525 Acc: 0.840
Epoch [88/150] Train Loss: 0.3165 Acc: 0.888	Val Loss: 0.4052 Acc: 0.860
Epoch [89/150] Train Loss: 0.3037 Acc: 0.899	Val Loss: 0.4449 Acc: 0.820
Epoch [90/150] Train Loss: 0.3171 Acc: 0.893	Val Loss: 0.3949 Acc: 0.840
Epoch [91/150] Train Loss: 0.3104 Acc: 0.900	Val Loss: 0.4210 Acc: 0.820
Epoch [92/150] Train Loss: 0.3266 Acc: 0.888	Val Loss: 0.4269 Acc: 0.840
Epoch [93/150] Train Loss: 0.3328 Acc: 0.883	Val Loss: 0.4197 Acc: 0.840
Epoch [94/150] Train Loss: 0.3184 Acc: 0.888	Val Loss: 0.3997 Acc: 0.840
Epoch [95/150] Train Loss: 0.3157 Acc: 0.886	Val Loss: 0.4032 Acc: 0.840
Epoch [96/150] Train Loss: 0.3153 Acc: 0.893	Val Loss: 0.3854 Acc: 0.840
Epoch [97/150] Train Loss: 0.3119 Acc: 0.892	Val Loss: 0.4006 Acc: 0.820
Epoch [98/150] Train Loss: 0.3180 Acc: 0.888	Val Loss: 0.4142 Acc: 0.840
Epoch [99/150] Train Loss: 0.3164 Acc: 0.893	Val Loss: 0.3982 Acc: 0.860
Epoch [100/150] Train Loss: 0.2996 Acc: 0.898	Val Loss: 0.3802 Acc: 0.840
Epoch [101/150] Train Loss: 0.3181 Acc: 0.897	Val Loss: 0.4166 Acc: 0.860
Epoch [102/150] Train Loss: 0.3134 Acc: 0.898	Val Loss: 0.4052 Acc: 0.820
Epoch [103/150] Train Loss: 0.3261 Acc: 0.894	Val Loss: 0.4462 Acc: 0.840
Epoch [104/150] Train Loss: 0.3112 Acc: 0.896	Val Loss: 0.4103 Acc: 0.880
Epoch [105/150] Train Loss: 0.3148 Acc: 0.895	Val Loss: 0.4056 Acc: 0.840
Epoch [106/150] Train Loss: 0.3207 Acc: 0.895	Val Loss: 0.4363 Acc: 0.860
Epoch [107/150] Train Loss: 0.3339 Acc: 0.883	Val Loss: 0.3964 Acc: 0.840
Epoch [108/150] Train Loss: 0.3150 Acc: 0.893	Val Loss: 0.4250 Acc: 0.840
Epoch [109/150] Train Loss: 0.3110 Acc: 0.899	Val Loss: 0.4186 Acc: 0.840
Epoch [110/150] Train Loss: 0.3086 Acc: 0.895	Val Loss: 0.4262 Acc: 0.820
Epoch [111/150] Train Loss: 0.3223 Acc: 0.889	Val Loss: 0.4288 Acc: 0.840
Epoch [112/150] Train Loss: 0.3259 Acc: 0.893	Val Loss: 0.3959 Acc: 0.840
Epoch [113/150] Train Loss: 0.3125 Acc: 0.893	Val Loss: 0.4321 Acc: 0.840
Epoch [114/150] Train Loss: 0.3147 Acc: 0.892	Val Loss: 0.4830 Acc: 0.820
Epoch [115/150] Train Loss: 0.3227 Acc: 0.890	Val Loss: 0.4215 Acc: 0.840
Epoch [116/150] Train Loss: 0.3200 Acc: 0.895	Val Loss: 0.4361 Acc: 0.860

Epoch [117/150] Train Loss: 0.3250 Acc: 0.891	Val Loss: 0.4109 Acc: 0.840
Epoch [118/150] Train Loss: 0.3063 Acc: 0.896	Val Loss: 0.4396 Acc: 0.860
Epoch [119/150] Train Loss: 0.3337 Acc: 0.881	Val Loss: 0.4502 Acc: 0.820
Epoch [120/150] Train Loss: 0.3163 Acc: 0.895	Val Loss: 0.4119 Acc: 0.840
Epoch [121/150] Train Loss: 0.3133 Acc: 0.895	Val Loss: 0.4085 Acc: 0.860
Epoch [122/150] Train Loss: 0.2998 Acc: 0.898	Val Loss: 0.4303 Acc: 0.820
Epoch [123/150] Train Loss: 0.3204 Acc: 0.890	Val Loss: 0.4479 Acc: 0.860
Epoch [124/150] Train Loss: 0.2993 Acc: 0.897	Val Loss: 0.4158 Acc: 0.840
Epoch [125/150] Train Loss: 0.3046 Acc: 0.903	Val Loss: 0.4217 Acc: 0.820
Epoch [126/150] Train Loss: 0.3095 Acc: 0.897	Val Loss: 0.4400 Acc: 0.840
Epoch [127/150] Train Loss: 0.3175 Acc: 0.888	Val Loss: 0.3932 Acc: 0.840
Epoch [128/150] Train Loss: 0.3173 Acc: 0.892	Val Loss: 0.4480 Acc: 0.820
Epoch [129/150] Train Loss: 0.3083 Acc: 0.899	Val Loss: 0.4057 Acc: 0.820
Epoch [130/150] Train Loss: 0.3289 Acc: 0.885	Val Loss: 0.4158 Acc: 0.840
Epoch [131/150] Train Loss: 0.3277 Acc: 0.886	Val Loss: 0.4119 Acc: 0.840
Epoch [132/150] Train Loss: 0.2986 Acc: 0.893	Val Loss: 0.3989 Acc: 0.860
Epoch [133/150] Train Loss: 0.3308 Acc: 0.894	Val Loss: 0.4284 Acc: 0.840
Epoch [134/150] Train Loss: 0.2944 Acc: 0.896	Val Loss: 0.4104 Acc: 0.840
Epoch [135/150] Train Loss: 0.3205 Acc: 0.894	Val Loss: 0.4324 Acc: 0.840
Epoch [136/150] Train Loss: 0.3025 Acc: 0.900	Val Loss: 0.4755 Acc: 0.800
Epoch [137/150] Train Loss: 0.3040 Acc: 0.897	Val Loss: 0.3841 Acc: 0.840
Epoch [138/150] Train Loss: 0.3040 Acc: 0.900	Val Loss: 0.4118 Acc: 0.840
Epoch [139/150] Train Loss: 0.3168 Acc: 0.890	Val Loss: 0.4268 Acc: 0.840
Epoch [140/150] Train Loss: 0.3133 Acc: 0.892	Val Loss: 0.4629 Acc: 0.840
Epoch [141/150] Train Loss: 0.3291 Acc: 0.890	Val Loss: 0.4331 Acc: 0.840
Epoch [142/150] Train Loss: 0.3254 Acc: 0.884	Val Loss: 0.4377 Acc: 0.840
Epoch [143/150] Train Loss: 0.3065 Acc: 0.898	Val Loss: 0.4269 Acc: 0.840
Epoch [144/150] Train Loss: 0.3271 Acc: 0.891	Val Loss: 0.4325 Acc: 0.840
Epoch [145/150] Train Loss: 0.2909 Acc: 0.906	Val Loss: 0.4253 Acc: 0.820
Epoch [146/150] Train Loss: 0.3242 Acc: 0.893	Val Loss: 0.4392 Acc: 0.840
Epoch [147/150] Train Loss: 0.3048 Acc: 0.900	Val Loss: 0.4444 Acc: 0.840
Epoch [148/150] Train Loss: 0.3026 Acc: 0.897	Val Loss: 0.4418 Acc: 0.840
Epoch [149/150] Train Loss: 0.3156 Acc: 0.896	Val Loss: 0.4188 Acc: 0.840
Epoch [150/150] Train Loss: 0.3209 Acc: 0.894	Val Loss: 0.4162 Acc: 0.840

Best validation accuracy: 0.880 at epoch 66

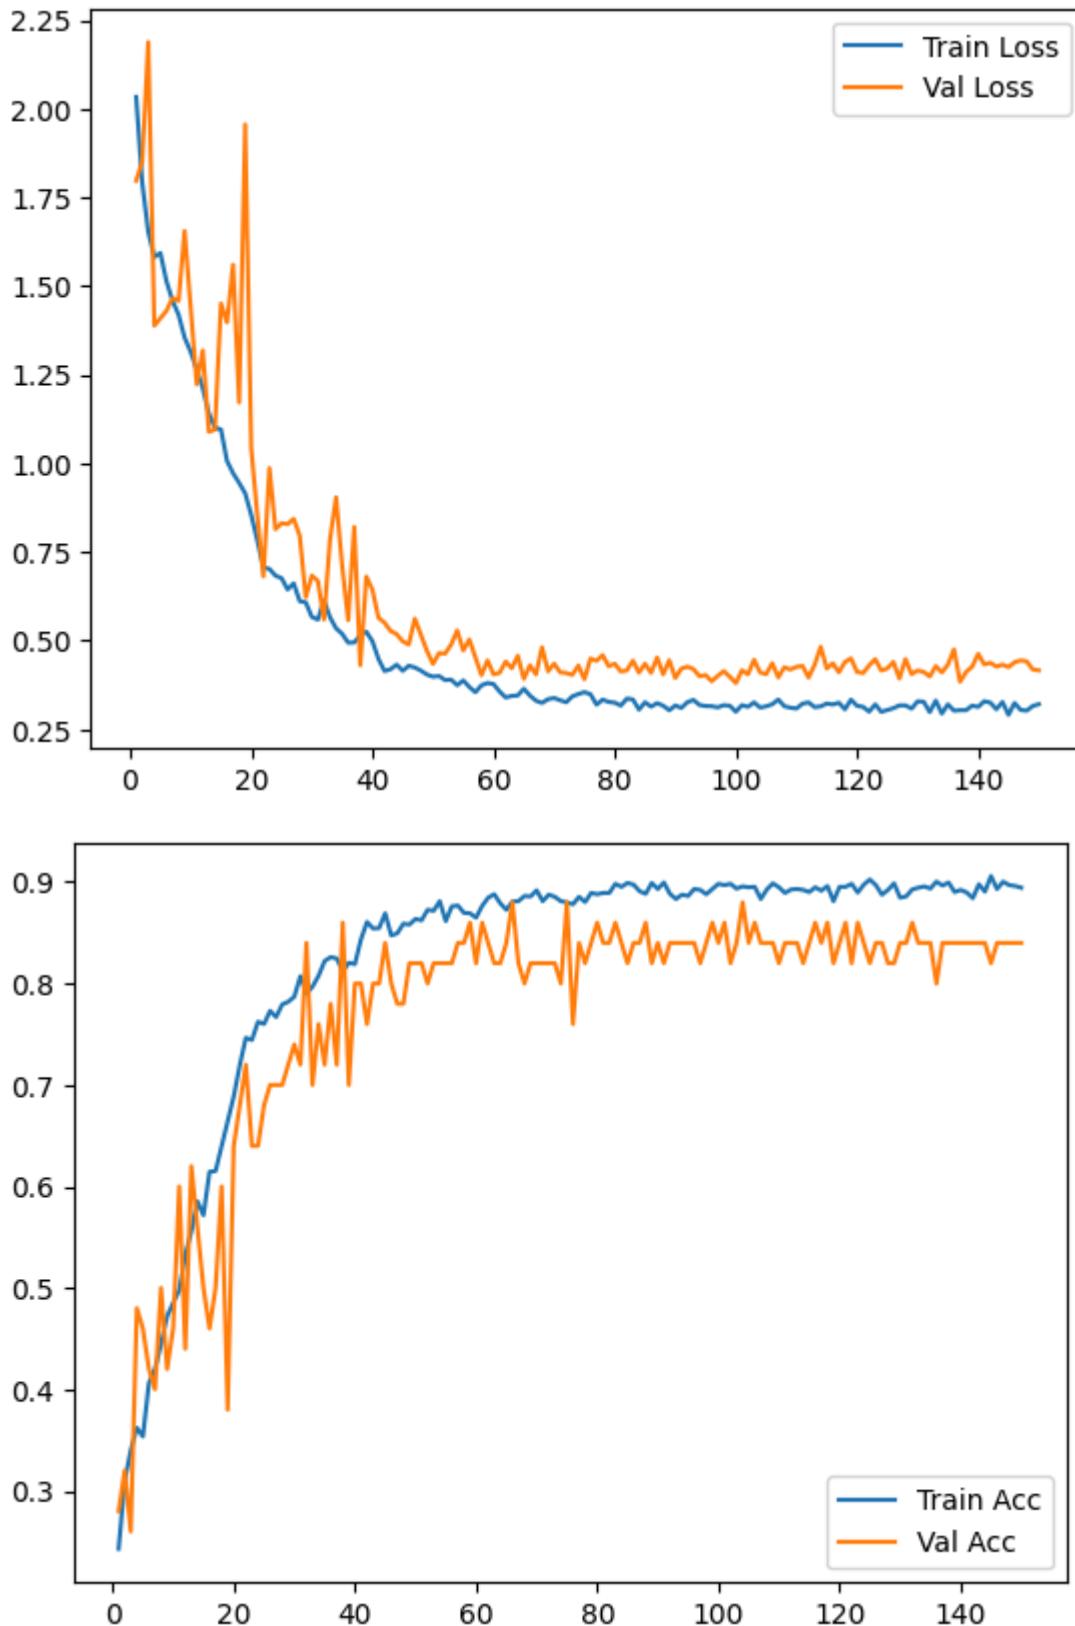
Test Accuracy: 0.860

Saved best model weights to: HW1_211644240_318801792_Q4.pkl

```
In [33]: epochs = range(1, len(train_loss_hist) + 1)

plt.figure()
plt.plot(epochs, train_loss_hist, label="Train Loss")
plt.plot(epochs, val_loss_hist, label="Val Loss")
plt.legend()
plt.show()

plt.figure()
plt.plot(epochs, train_acc_hist, label="Train Acc")
plt.plot(epochs, val_acc_hist, label="Val Acc")
plt.legend()
plt.show()
```



II. Analyzing a Pre-trained CNN (Filters) (10pt)

In this part, you are going to analyze a (large) pre-trained model. Pre-trained models are quite popular these days, as big companies can train really large models on large datasets (something that personal users can't do as they lack the sufficient hardware). These pre-trained models can be used to fine-tune on other/small datasets or used as components in other tasks (like using a pre-trained classifier for object detection).

All pre-trained models expect input images normalized in the same way, i.e. mini-batches of 3-channel RGB images of shape (3 x H x W), where H and W are expected to be at least 224. The images have to be loaded in to a range of [0, 1] and then normalized using mean = [0.485, 0.456, 0.406] and std = [0.229, 0.224, 0.225].

You can use the following transform to normalize:

```
normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]) Read more here
```

1. Load a pre-trained VGG16 with PyTorch using
`torchvision.models.vgg16(pretrained=True, progress=True, **kwargs)`
[\(read more here\)](#). Don't forget to use the model in evaluation mode
`(model.eval())`.
2. Load the images in the `hw1_data/birds` folder and display them.
3. Pre-process the images to fit VGG16's architecture. What steps did you take?
4. Feed the images (forward pass) to the model. What are the outputs?
5. Choose an image of a dog in the `hw1_data/dogs` folder, display it and feed it to network. What are the outputs?
6. For the first 3 filters in the first layer of VGG16, plot their response (their output) for the image from section 5. Explain what do you see.

1. Load a pre-trained VGG16

`VGG16_Weights.DEFAULT` loads ImageNet pre-trained weights and also provides metadata such as the list of 1000 ImageNet class names.

```
In [27]: import torch
from torchvision import transforms
from torchvision.models import vgg16, VGG16_Weights
from PIL import Image
from pathlib import Path
import matplotlib.pyplot as plt

%matplotlib inline

weights = VGG16_Weights.DEFAULT
model = vgg16(weights=weights)
model.eval()
imagenet_classes = weights.meta["categories"]
```

2. Load the images in the `birds` folder and display them

The project directory contains a folder `birds/` with files. We read all images from that folder and show them.

```
In [28]: birds_dir = Path("birds")

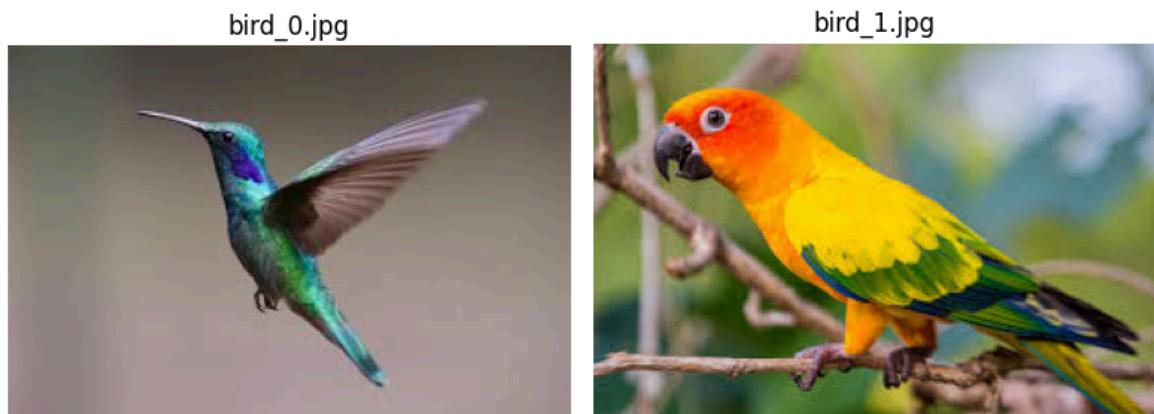
bird_paths = sorted([
    p for p in birds_dir.iterdir()
    if p.suffix.lower() in [".jpg", ".jpeg", ".png"]
])

fig, axes = plt.subplots(1, len(bird_paths), figsize=(4 * len(bird_paths), 4))

if len(bird_paths) == 1:
    axes = [axes]

for ax, img_path in zip(axes, bird_paths):
    img = Image.open(img_path).convert("RGB")
    ax.imshow(img)
    ax.set_title(img_path.name)
    ax.axis("off")

plt.tight_layout()
plt.show()
```



3. Pre-process the images to fit VGG16

VGG16 expects:

- 3-channel RGB images,
- of size at least 224×224 ,
- with pixel values in the range $[0, 1]$,
- normalized channel-wise using:
 - mean = $[0.485, 0.456, 0.406]$
 - std = $[0.229, 0.224, 0.225]$

To achieve this, we apply the following transformation of steps:

1. `Resize(256)` - rescale the image so that the shorter side is 256 pixels.
2. `CenterCrop(224)` - take a 224×224 crop from the center.
3. `ToTensor()` - convert the image to a PyTorch tensor of shape $(3, H, W)$ with values in $[0, 1]$.

4. `Normalize(mean, std)` - normalize each channel using the given mean and std.

We then stack all processed images into a mini-batch of shape (N, 3, 224, 224), where N is the number of bird images which is 2 in this particular file.

```
In [29]: # Define preprocessing transform as defined in the question
normalize = transforms.Normalize(
    mean=[0.485, 0.456, 0.406],
    std=[0.229, 0.224, 0.225]
)

preprocess = transforms.Compose([
    transforms.Resize(256),           # 1. resize short side to 256
    transforms.CenterCrop(224),        # 2. center crop to 224x224
    transforms.ToTensor(),            # 3. convert to tensor [0,1]
    normalize                         # 4. normalize with ImageNet stats
])

bird_tensors = []

for img_path in bird_paths:
    img = Image.open(img_path).convert("RGB")
    x = preprocess(img)
    bird_tensors.append(x)

# Stack into batch: (N, 3, 224, 224)
bird_batch = torch.stack(bird_tensors, dim=0)
```

4. Forward pass for the bird images - outputs

We perform a forward pass through VGG16 using the bird batch.

The raw network output for each image is a 1000-dimensional vector of **logits**, one score per ImageNet class.

We then apply `softmax` to obtain a probability distribution over the 1000 classes, and print the top-5 predictions for each bird image.

```
In [30]: with torch.no_grad():
    outputs_birds = model(bird_batch) # shape: (N, 1000)

print("Raw outputs shape (logits):", outputs_birds.shape)

# Convert logits to probabilities
probs_birds = outputs_birds.softmax(dim=1)

top5_prob, top5_idx = probs_birds.topk(5, dim=1)

for i, img_path in enumerate(bird_paths):
    print(f"\nImage: {img_path.name}")
    for rank in range(5):
        idx = top5_idx[i, rank].item()
        prob = top5_prob[i, rank].item()
```

```
label = imagenet_classes[idx]
print(f" #{{rank+1}}: {{label:30s}} prob={{prob:.4f}}")
```

Raw outputs shape (logits): torch.Size([2, 1000])

```
Image: bird_0.jpg
#1: hummingbird           prob=0.8644
#2: jacamar               prob=0.1355
#3: bee eater              prob=0.0002
#4: black stork             prob=0.0000
#5: water ouzel             prob=0.0000

Image: bird_1.jpg
#1: lorikeet                prob=0.9415
#2: macaw                   prob=0.0401
#3: toucan                  prob=0.0108
#4: bee eater                 prob=0.0029
#5: goldfinch                prob=0.0017
```

5. Dog image: display and predictions

Now we choose one image from the `dogs/` folder, apply the same pre-processing, and run it through the network. Again, we inspect the top-5 predicted classes.

We expect the model to output different dog breeds as its most probable classes.

```
In [31]: dogs_dir = Path("dogs")
dog_paths = sorted([
    p for p in dogs_dir.iterdir()
    if p.suffix.lower() in [".jpg", ".jpeg", ".png"]
])

# Pick one dog image
dog_path = dog_paths[0]
dog_img = Image.open(dog_path).convert("RGB")
plt.imshow(dog_img)
plt.title(dog_path.name)
plt.axis("off")
plt.show()

# Preprocess and create batch of size 1
dog_tensor = preprocess(dog_img).unsqueeze(0) # shape: (1, 3, 224, 224)

# Forward pass
with torch.no_grad():
    outputs_dog = model(dog_tensor) # shape: (1, 1000)

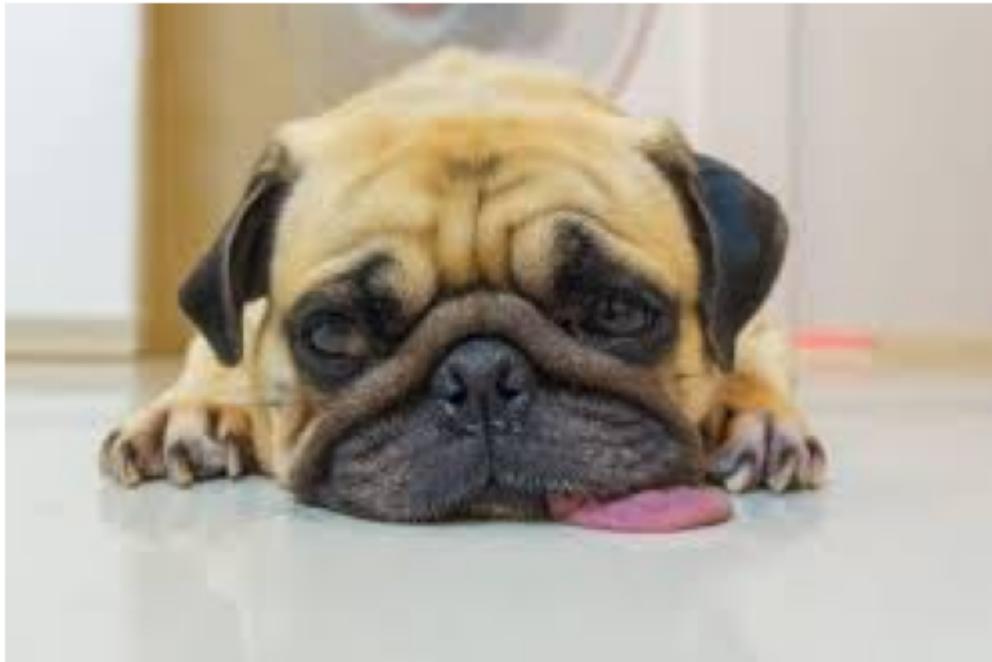
print("Dog output shape (logits):", outputs_dog.shape)

# Top-5 predictions
probs_dog = outputs_dog.softmax(dim=1)
top5_prob_dog, top5_idx_dog = probs_dog.topk(5, dim=1)

print(f"\nPredictions for {dog_path.name}:")
for rank in range(5):
    idx = top5_idx_dog[0, rank].item()
    prob = top5_prob_dog[0, rank].item()
```

```
label = imagenet_classes[idx]
print(f" # {rank+1}: {label:30s} prob={prob:.4f}")
```

dog_0.jpg



Dog output shape (logits): torch.Size([1, 1000])

Predictions for dog_0.jpg:

#1: pug	prob=0.9856
#2: Brabancon griffon	prob=0.0054
#3: bull mastiff	prob=0.0041
#4: French bulldog	prob=0.0040
#5: boxer	prob=0.0006

6. First-layer filter responses for the dog image

We now inspect the first convolutional layer of VGG16 (`features[0]`). It is a `Conv2d` layer with 3 input channels and 64 output channels (64 filters).

We:

1. Feed the pre-processed dog image through this single convolutional layer.
2. Take the output feature maps of the first 3 filters (channels 0, 1, 2).
3. Normalize each feature map for visualization and plot them as grayscale images.

These feature maps show where each filter "fires" on the image.

```
In [32]: # First conv layer of VGG16
first_conv = model.features[0]
print(first_conv)

with torch.no_grad():
    first_layer_output = first_conv(dog_tensor)

activations = first_layer_output[0]
fig, axes = plt.subplots(1, 3, figsize=(12, 4))
```

```

for i in range(3):
    act = activations[i].cpu().numpy()

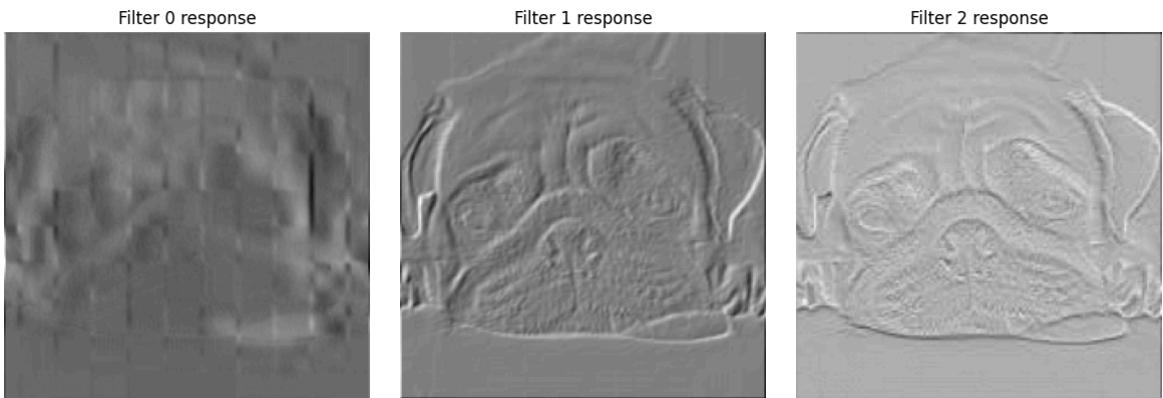
    # Normalize each feature map to [0,1] for visualization
    act_min, act_max = act.min(), act.max()
    act_vis = (act - act_min) / (act_max - act_min + 1e-8)

    axes[i].imshow(act_vis, cmap="gray")
    axes[i].set_title(f"Filter {i} response")
    axes[i].axis("off")

plt.tight_layout()
plt.show()

```

Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))



Explanation of what we see

Each grayscale image corresponds to the activation map of one convolution filter for the dog image.

- Bright regions indicate locations where the filter produces a high response, meaning the local image pattern matches the filter's preferred pattern.
- Dark regions indicate low response.

Because this is the first convolutional layer, its filters typically learn very low-level features such as:

- oriented edges (vertical, horizontal, diagonal),
- simple color or intensity contrasts,
- small texture patterns.

We can see that different filters highlight different structures of the dog image: for example, one filter may respond strongly along the dog's outline, another may focus on fur texture, and another may emphasize contrast between the dog and the background.

These simple feature maps are then combined and transformed by deeper layers to build up increasingly complex and abstract representations.

In []: